The Siri we should have. An AI-native phone OS, built security-first.

What it is. You say it, agents do it in your real apps. The model proposes; a gate in the OS decides; risky steps wait for a tap or your fingerprint; payments and passwords always require a fingerprint; everything is logged in a hash chain.

Where it stands. Working on a rooted Moto G 5G (2023), Android 13, our lab phone. Before this week: an island overlay prototype, SystemUI experiments, Codex running on the phone. Built Oct 26 to 28, 123 commits: the gate (leases, Rule of Two, on-device Prompt Guard 2, fingerprint-gated approvals, audit chain, per-app kernel firewall); agentd with Codex, Claude or an on-device model; cross-app memory; proactive cards; scheduled tasks on a locked phone; automatic offline fallback; live voice; an AI keyboard; app mods and micro-apps; the launcher.

Company. Utah-based founder, Utah LLC already formed: eligible now for UTIF microgrants, SBIR/STTR with Nucleus Grow, and the Nucleus Fund. **90-day goal.** From one rooted lab phone to a shippable build, a 10-person Utah pilot, and an SBIR Phase I proposal ready to submit.

~30 → ~11 s
median task time this week, real apps · agentd/bench/RESULTS.md

26/30
real tasks driven by the on-device 0.8B model, was 0/30 · runs/2026-10-27-local-model.md

38/38 · 12/12
cross-app memory, tuning and held-out · memory/eval/RESULTS.md

7.3 → 3.5 s
keyboard suggestions, typing pause · keyboard/LATENCY.md

MILESTONES

WHEN	MILESTONE	DONE WHEN
Days 1–30	Off root: managed-device (device-owner) gate, plus Android AppFunctions where apps support them. Open threat model and test suite (7 attack types from arXiv 2607.00333, plus our guard and memory evals). SBIR topic with Nucleus Grow; SAM.gov and SBA registrations; UTIF microgrant application.	Gate, audit and firewall on an unrooted phone. Public repo with results. Topic chosen, filings sent.
Days 31–60	10 Utah pilot users through a target channel such as the Utah Assistive Technology Program (USU) or an independent living center, plain-language consent. Independent red-team review of the gate.	10 signed consents, phones set up. Written report, fixes merged.
Days 61–90	Run the pilot: task success, approvals per task, attacks blocked, time saved. SBIR Phase I proposal (NSF, NIH/NIA or DoD). Apply to the Utah Office of AI Policy Learning Lab. Results to the Nucleus Fund.	Pilot report published. Proposal ready for the first open solicitation.

SBIR timing depends on agencies reopening solicitations after the April 2026 reauthorization, so "ready to submit" is the honest day-90 target.

USE OF THE $10K



Devices 6 low-cost Android phones (~$150–250) + 2 Pixels: off root, onto the shippable path

Utah pilot $150 stipends × 10, assistive-tech specialist time, travel to Logan / Uintah Basin

Red team independent review against published mobile-agent attacks

Credits cloud + local models

Legal SBIR filings, consent forms

Beyond money: introductions to UATP (USU), DSPD, a university co-PI, and Nucleus Grow.

PEOPLE

Founder (Josh)	Product, security design, reviews every security decision.	In place
AI coding agents	Claude Code, Codex, Fable design agents: build, test, research under the founder's review.	In place
University co-PI	U of U Kahlert School of Computing or USU IDRPP; required for STTR.	Needed, intro
AT specialist	Runs the pilot safely with disabled and older users.	Pilot budget
Android engineer	Part-time: managed-device and custom-build paths.	After this bounty
Security reviewer	Independent proof of the safety claim.	Red-team budget

TOP RISKS

Root isn't shippable	Root is the lab. Ship as a managed-device app, a privileged app on a re-locked build, or via AppFunctions.
Google or Apple ships it	Ours is model-agnostic, runs on cheap phones, keeps working offline, and the user or agency owns the rules and the log.
Pilot safety	No send, pay or delete without a person's approval on that exact action. Attack demos use only our own test apps.
Provider terms	Product uses API keys or local models, not personal subscriptions.

Sources: every number is from a results file in the repo; per-claim sources in pitch/claims-check.md and pitch/deck/notes.md. Voice: 1.7 s from end of speech to task start, spoken "done" ~0.9 s after (spikes/voice/results.jsonl, 7 synthetic-voice runs). "Target channel" means no agreement signed yet. pitch/one-pager.md